

- ✓ Title: Data Efficiency Matters: Performance from 10% → 100% Training Size

Models: Linear Regression, Decision Tree, and Neural Network (MLP) bold text

Dataset: UCI Abalone

```
# ## Step 1: Import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

import warnings
warnings.filterwarnings('ignore')
```

```
# ## Step 2: Load Dataset

url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/abalone/abalone.data'
columns = ['Sex', 'Length', 'Diameter', 'Height', 'WholeWeight', 'ShuckedWeight', 'VisceraWeight', 'ShellWeight', 'Rings']
df = pd.read_csv(url, names=columns)
```

```
print("Dataset shape:", df.shape)
df.head()
```

Dataset shape: (4177, 9)									
	Sex	Length	Diameter	Height	WholeWeight	ShuckedWeight	VisceraWeight	ShellWeight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
# ## Step 3: Basic Information & Descriptive Statistics
```

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   Sex          4177 non-null   object  
 1   Length       4177 non-null   float64 
 2   Diameter     4177 non-null   float64 
 3   Height       4177 non-null   float64 
 4   WholeWeight  4177 non-null   float64 
 5   ShuckedWeight 4177 non-null   float64 
 6   VisceraWeight 4177 non-null   float64 
 7   ShellWeight  4177 non-null   float64 
 8   Rings         4177 non-null   int64  
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
None
```

```
print(df.describe())
```

	Length	Diameter	Height	WholeWeight	ShuckedWeight	VisceraWeight	ShellWeight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367			
std	0.120093	0.099240	0.041827	0.490389	0.221963			
min	0.075000	0.055000	0.000000	0.002000	0.001000			
25%	0.450000	0.350000	0.115000	0.441500	0.186000			
50%	0.545000	0.425000	0.140000	0.799500	0.336000			
75%	0.615000	0.480000	0.165000	1.153000	0.502000			
max	0.815000	0.650000	1.130000	2.825500	1.488000			

mean	0.180594	0.238831	9.933684
std	0.109614	0.139203	3.224169
min	0.000500	0.001500	1.000000
25%	0.093500	0.130000	8.000000
50%	0.171000	0.234000	9.000000
75%	0.253000	0.329000	11.000000
max	0.760000	1.005000	29.000000

```
print("\nMissing Values:\n", df.isnull().sum())
```

```
Missing Values:
Sex           0
Length        0
Diameter      0
Height         0
WholeWeight    0
ShuckedWeight  0
VisceraWeight  0
ShellWeight    0
Rings          0
dtype: int64
```

```
# ## Step 4: Handle Categorical Feature & Scaling
```

```
df = pd.get_dummies(df, columns=['Sex'], drop_first=True)
X = df.drop('Rings', axis=1)
y = df['Rings']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
# ## Step 5: Outlier Detection (IQR Method)
```

```
print("\nOutlier Detection using IQR Method:")

# Select only numeric columns for IQR
numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns

# Calculate IQR
Q1 = df[numeric_cols].quantile(0.25)
Q3 = df[numeric_cols].quantile(0.75)
IQR = Q3 - Q1
#Define bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR


# Keep only rows within bounds
df_clean = df[~((df[numeric_cols] < lower_bound) | (df[numeric_cols] > upper_bound)).any(axis=1)]

print(f"Original shape: {df.shape}")
print(f"After removing outliers: {df_clean.shape}")
print(f"Rows removed: {df.shape[0] - df_clean.shape[0]}")
```


```
Outlier Detection using IQR Method:
Original shape: (4177, 10)
After removing outliers: (3781, 10)
Rows removed: 396
```

```
# ## Step 6: Data Visualization
```

```
# Correlation Heatmap
plt.figure(figsize=(10,6))
sns.heatmap(df_clean.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Feature Correlation Heatmap (Abalone Dataset)")
plt.show()
```




```
plt.figure(figsize=(6,4))
sns.histplot(df['Rings'], kde=True, bins=20)
plt.title('Target Distribution (Rings)')
plt.show()
```




```
# Pairplot for a few important columns
plt.figure(figsize=(6,4))
sns.pairplot(df_clean[['Length', 'Diameter', 'Height', 'ShellWeight', 'Rings']])
plt.suptitle("Pairplot: Relationships between Key Features", y=1.02)
plt.show()
```

<Figure size 600x400 with 0 Axes>


Pairplot: Relationships between Key Features



```
# Boxplots for numeric columns
plt.figure(figsize=(6,4))
df_clean.boxplot(rot=45)
plt.title("Boxplots of Numeric Features after Outlier Removal")
plt.show()
```




```
plt.figure(figsize=(8,5))
sns.histplot(df_clean['Rings'], bins=15, kde=True, color='purple')
plt.title("Distribution of Target Variable: Rings")
plt.xlabel("Rings (Age Indicator)")
plt.ylabel("Frequency")
plt.show()
```



```
corr_target = df_clean.corr()['Rings'].sort_values(ascending=False)
plt.figure(figsize=(8,4))
sns.barplot(x=corr_target.index, y=corr_target.values, palette='viridis')
plt.title("Feature Correlation Strength with Target (Rings)")
plt.xticks(rotation=45)
plt.show()
```

Feature Correlation Strength with Target (Rings)



Step 7: Train-Test Split

```
X = df.drop('Rings', axis=1)
y = df['Rings']
X_train_pool, X_test, y_train_pool, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print("Train size:", X_train_pool.shape, "Test size:", X_test.shape)

Train size: (3341, 9) Test size: (836, 9)
```

```
# ## Step 8: Function to Evaluate Models
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

def evaluate_model(y_true, y_pred):
    # Compute metrics manually to avoid version issues
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse) # manually compute root MSE
    mae = mean_absolute_error(y_true, y_pred)
    r2 = r2_score(y_true, y_pred)
    return rmse, mae, r2
```

```
# Initialize models
lr = LinearRegression()
dt = DecisionTreeRegressor(random_state=42)
nn = MLPRegressor(hidden_layer_sizes=(64, 32), max_iter=1000, random_state=42)
```

```
#  Step 9: Train models on different training sizes and collect results
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

train_fractions = [0.1, 0.3, 0.5, 0.7, 1.0]

results_lr = []
results_dt = []
results_nn = []

for frac in train_fractions:
    #  Handle 100% case separately (no split)
    if frac < 1.0:
        X_train_frac, _, y_train_frac, _ = train_test_split(X_train, y_train, train_size=frac, random_state=42)
    else:
        X_train_frac, y_train_frac = X_train, y_train

    # --- Linear Regression ---
    lr.fit(X_train_frac, y_train_frac)
    y_pred_lr = lr.predict(X_test)
    rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))
    r2_lr = r2_score(y_test, y_pred_lr)
    results_lr.append((frac, rmse_lr, len(X_train_frac), r2_lr))

    # --- Decision Tree ---
    dt.fit(X_train_frac, y_train_frac)
```

```

y_pred_dt = dt.predict(X_test)
rmse_dt = np.sqrt(mean_squared_error(y_test, y_pred_dt))
r2_dt = r2_score(y_test, y_pred_dt)
results_dt.append((frac, rmse_dt, len(X_train_frac), r2_dt))

# --- Neural Network ---
nn.fit(X_train_frac, y_train_frac)
y_pred_nn = nn.predict(X_test)
rmse_nn = np.sqrt(mean_squared_error(y_test, y_pred_nn))
r2_nn = r2_score(y_test, y_pred_nn)
results_nn.append((frac, rmse_nn, len(X_train_frac), r2_nn))

# ✅ Print summary
print("✅ Linear Regression Results:", results_lr)
print("✅ Decision Tree Results:", results_dt)
print("✅ Neural Network Results:", results_nn)

```

```

.7, np.float64(2.2298345191197004), 2338, 0.5406867892879128), (1.0, np.float64(2.211613087121836), 3341, 0.5481628137889265)]  

np.float64(3.129386873008343), 2338, 0.09534668059620688), (1.0, np.float64(3.043936635219992), 3341, 0.14407663220938294)]  

0.7, np.float64(2.151770442364263), 2338, 0.5722839541083458), (1.0, np.float64(2.2058881038283484), 3341, 0.5504990382511511)]

```

```

# Convert results to DataFrames
df_lr = pd.DataFrame(results_lr, columns=['Training Fraction', 'RMSE', 'Samples Used', 'R²'])
df_dt = pd.DataFrame(results_dt, columns=['Training Fraction', 'RMSE', 'Samples Used', 'R²'])
df_nn = pd.DataFrame(results_nn, columns=['Training Fraction', 'RMSE', 'Samples Used', 'R²'])

# Display tables clearly
print("📊 Linear Regression Performance")
display(df_lr.style.background_gradient(cmap='Blues'))

print("\n🌳 Decision Tree Performance")
display(df_dt.style.background_gradient(cmap='Greens'))

print("\n🧠 Neural Network Performance")
display(df_nn.style.background_gradient(cmap='Reds'))

```

📊 Linear Regression Performance				
	Training Fraction	RMSE	Samples Used	R²
0	0.100000	2.274272	334	0.522197
1	0.300000	2.257619	1002	0.529169
2	0.500000	2.244504	1670	0.534624
3	0.700000	2.229835	2338	0.540687
4	1.000000	2.211613	3341	0.548163

🌳 Decision Tree Performance				
	Training Fraction	RMSE	Samples Used	R²
0	0.100000	3.093248	334	0.116120
1	0.300000	2.987614	1002	0.175458
2	0.500000	2.989215	1670	0.174574
3	0.700000	3.129387	2338	0.095347
4	1.000000	3.043937	3341	0.144077

🧠 Neural Network Performance				
	Training Fraction	RMSE	Samples Used	R²
0	0.100000	2.221576	334	0.544083
1	0.300000	2.289673	1002	0.515704
2	0.500000	2.190666	1670	0.556681
3	0.700000	2.151770	2338	0.572284
4	1.000000	2.205888	3341	0.550499

```

# Step 10: Visualize Model Performance vs Training Size

# Convert results into a DataFrame for easy plotting
train_df = pd.DataFrame(results, columns=['Train Size', 'RMSE', 'MAE', 'R²'])


# --- Plot R² Score vs Training Size ---
plt.figure(figsize=(7,5))

```


```
plt.plot(train_df['Train Size'], train_df['R2'], marker='o', color='green')
plt.title('Model R2 Score vs Training Data Size')
plt.xlabel('Training Data Fraction')
plt.ylabel('R2 Score')
plt.grid(True)
plt.show()
```



```
# --- Plot RMSE vs Training Size ---
plt.figure(figsize=(7,5))
plt.plot(train_df['Train Size'], train_df['RMSE'], marker='o', color='red')
plt.title('RMSE vs Training Data Size')
plt.xlabel('Training Data Fraction')
plt.ylabel('RMSE')
plt.grid(True)
plt.show()
```



```
# --- Plot MAE vs Training Size ---
plt.figure(figsize=(7,5))
plt.plot(train_df['Train Size'], train_df['MAE'], marker='o', color='orange')
plt.title('MAE vs Training Data Size')
plt.xlabel('Training Data Fraction')
plt.ylabel('MAE')
plt.grid(True)
plt.show()
```



Step 10B: Performance Comparison Visualization (RMSE & R² vs Training Size)

```
import matplotlib.pyplot as plt
import numpy as np

# Ensure results_lr, results_dt, and results_nn exist from Step 9

# Convert results into lists
train_sizes = [r[0]*100 for r in results_lr] # Convert 0.1 → 10%

rmse_lr = [r[1] for r in results_lr]
rmse_dt = [r[1] for r in results_dt]
rmse_nn = [r[1] for r in results_nn]

r2_lr = [r[3] if not np.isnan(r[3]) else 0 for r in results_lr]
r2_dt = [r[3] if not np.isnan(r[3]) else 0 for r in results_dt]
r2_nn = [r[3] if not np.isnan(r[3]) else 0 for r in results_nn]

# --- Plot RMSE comparison ---
plt.figure(figsize=(7,5))
plt.plot(train_sizes, rmse_lr, marker='o', label='Linear Regression', color='blue')
plt.plot(train_sizes, rmse_dt, marker='s', label='Decision Tree', color='green')
plt.plot(train_sizes, rmse_nn, marker='^', label='Neural Network', color='red')
plt.title("Model Performance (RMSE) vs Training Data Size", fontsize=13)
plt.xlabel("Training Size (%)")
plt.ylabel("Root Mean Squared Error (RMSE)")
plt.legend()
plt.grid(True)
plt.show()

# --- Plot R2 comparison ---
plt.figure(figsize=(7,5))
plt.plot(train_sizes, r2_lr, marker='o', label='Linear Regression', color='blue')
plt.plot(train_sizes, r2_dt, marker='s', label='Decision Tree', color='green')
plt.plot(train_sizes, r2_nn, marker='^', label='Neural Network', color='red')
plt.title("Model Performance (R2) vs Training Data Size", fontsize=13)
plt.xlabel("Training Size (%)")
plt.ylabel("R2 Score")
plt.legend()
plt.grid(True)
plt.show()

print("✅ Visualization completed – Models compared successfully!")
```

Model Performance (RMSE) vs Training Data Size

